

Touray_FinalProject

Sheikh-Sedat Touray

2023-10-15

```
heart <- read.csv('heart.csv')
head(heart)
```

```
##   age sex cp trtbps chol fbs restecg thalachh exng oldpeak slp caa thall output
## 1  63  1  3   145  233   1       0     150    0    2.3   0  0    1      1
## 2  37  1  2   130  250   0       1     187    0    3.5   0  0    2      1
## 3  41  0  1   130  204   0       0     172    0    1.4   2  0    2      1
## 4  56  1  1   120  236   0       1     178    0    0.8   2  0    2      1
## 5  57  0  0   120  354   0       1     163    1    0.6   2  0    2      1
## 6  57  1  0   140  192   0       1     148    0    0.4   1  0    1      1
```

Ckeck for missing and Null Values

```
nullv <- is.null(heart)
head(nullv)
```

```
## [1] FALSE
```

```
missnah <- is.na(heart)
head(missnah)
```

```
##      age  sex  cp trtbps  chol  fbs restecg thalachh  exng oldpeak  slp
## [1,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
## [2,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
## [3,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
## [4,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
## [5,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
## [6,] FALSE FALSE FALSE  FALSE FALSE FALSE  FALSE  FALSE FALSE  FALSE FALSE
##      caa thall output
## [1,] FALSE FALSE  FALSE
## [2,] FALSE FALSE  FALSE
## [3,] FALSE FALSE  FALSE
## [4,] FALSE FALSE  FALSE
## [5,] FALSE FALSE  FALSE
## [6,] FALSE FALSE  FALSE
```


```
library(plyr)
heart$output <- as.factor(heart$output)
heart$output <- revalue(heart$output, c('1' = 'Yes', '0' = 'No'))
heart$sex <- as.factor(heart$sex)
heart$sex <- revalue(heart$sex, c('1' = 'Male', '0' = 'Female'))
heart$cp <- as.factor(heart$cp)
heart$cp <- revalue(heart$cp, c('0' = 'typical', '1' = 'atypical', '2' = 'non-aginal', '3'='Asymptomatic'))
heart$restecg <- as.factor(heart$restecg)
heart$restecg <- revalue(heart$restecg, c('0' = 'Normal', '1' = 'Abnormal', '2' = 'LVH'))
```

```

heart$exng <- as.factor(heart$exng)
heart$exng <- revalue(heart$exng, c('1' = 'Yes', '0' = 'No'))
heart$slp <- as.factor(heart$slp)
heart$slp <- revalue(heart$slp, c('0' = 'Upsloping', '1' = 'Flat', '2' = 'Downsloping'))
heart$fbs <- as.factor(heart$fbs)
heart$fbs <- revalue(heart$fbs, c('1' = 'True', '0' = 'False'))

library(ggplot2)
par(mfrow=c(3,2))
classofpar <- ggplot(heart, aes(x=output, fill=output))+
  geom_bar()+
  xlab("Heart Attack")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack")
classofpar

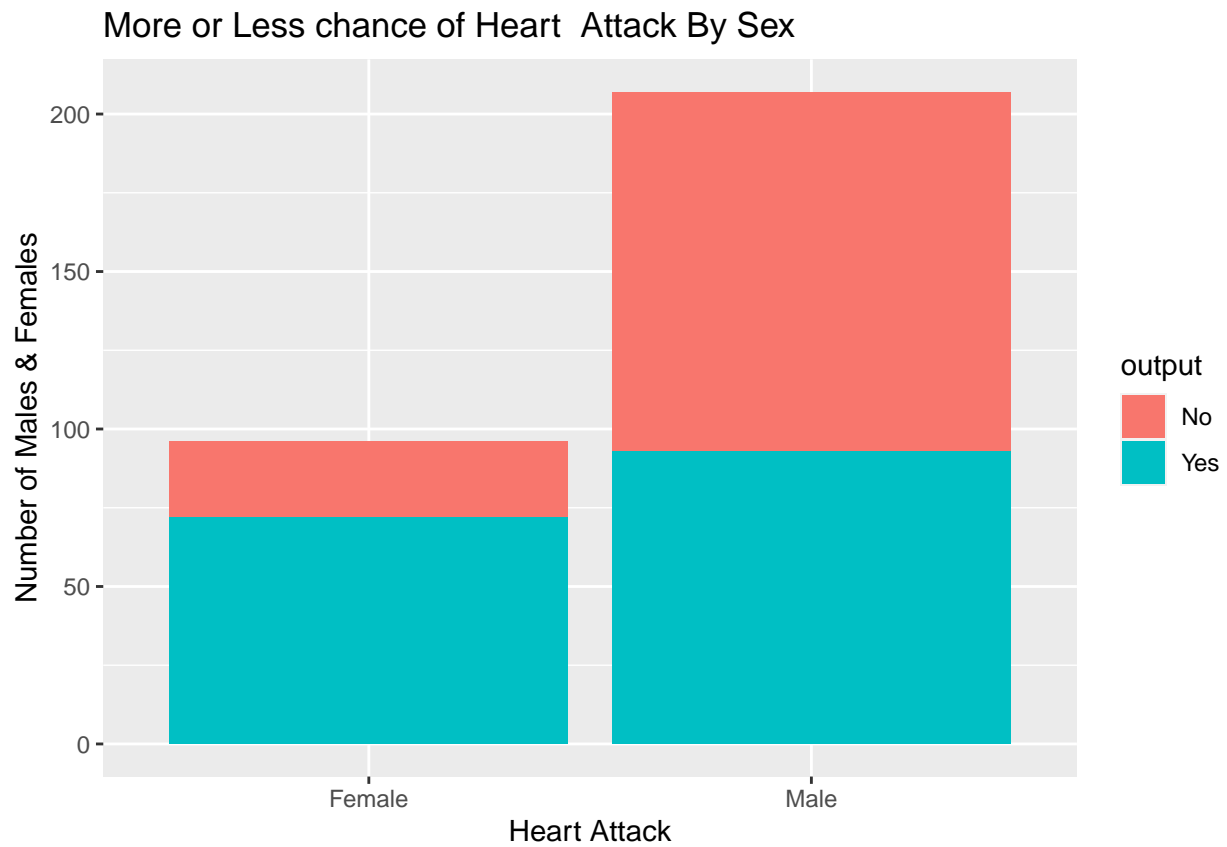
```




```

classofpar1 <- ggplot(heart, aes(x=sex, fill=output))+
  geom_bar()+
  xlab("Heart Attack")+
  ylab("Number of Males & Females")+
  ggtitle("More or Less chance of Heart Attack By Sex")
classofpar1


```



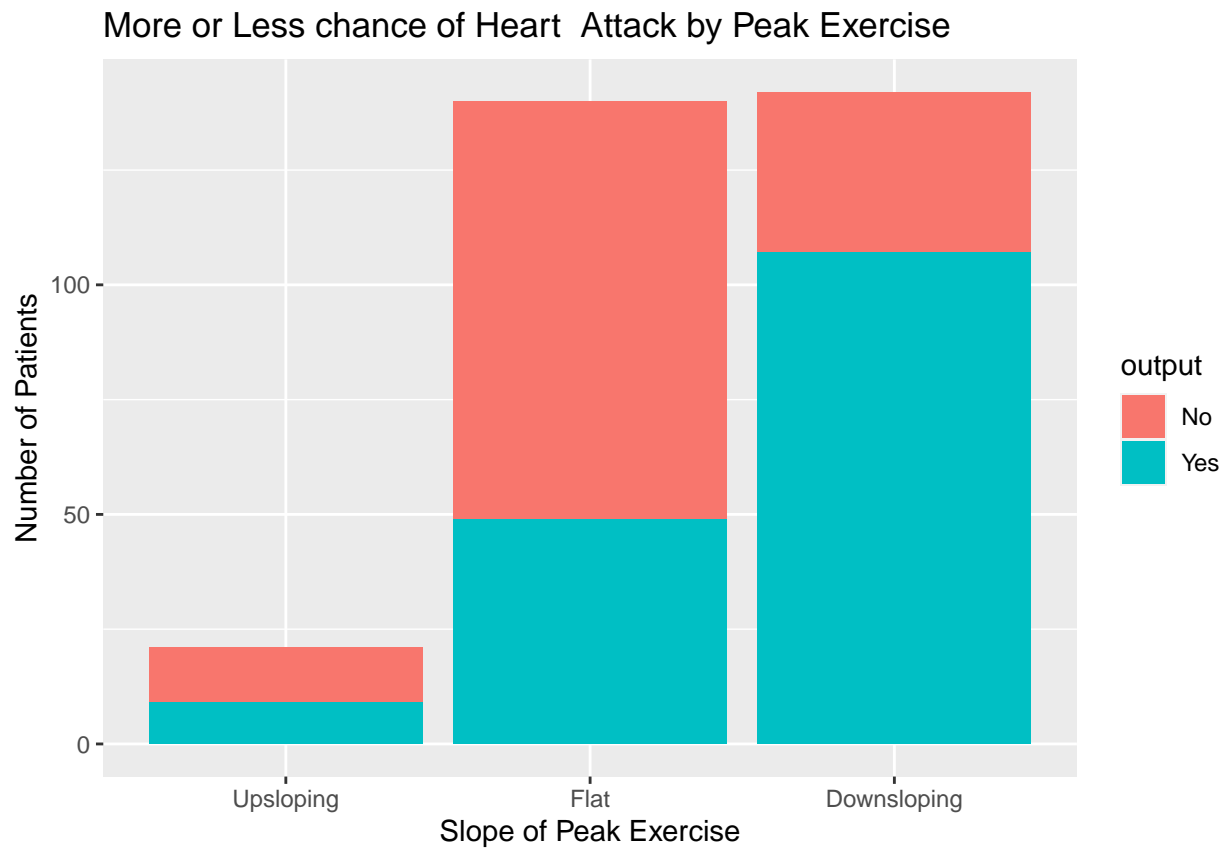
```
library(ggplot2)
classofpar2 <- ggplot(heart, aes(x=cp, fill=output))+
  geom_bar()+
  xlab("Type of chest pain")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack by Chest Pain")
classofpar2
```



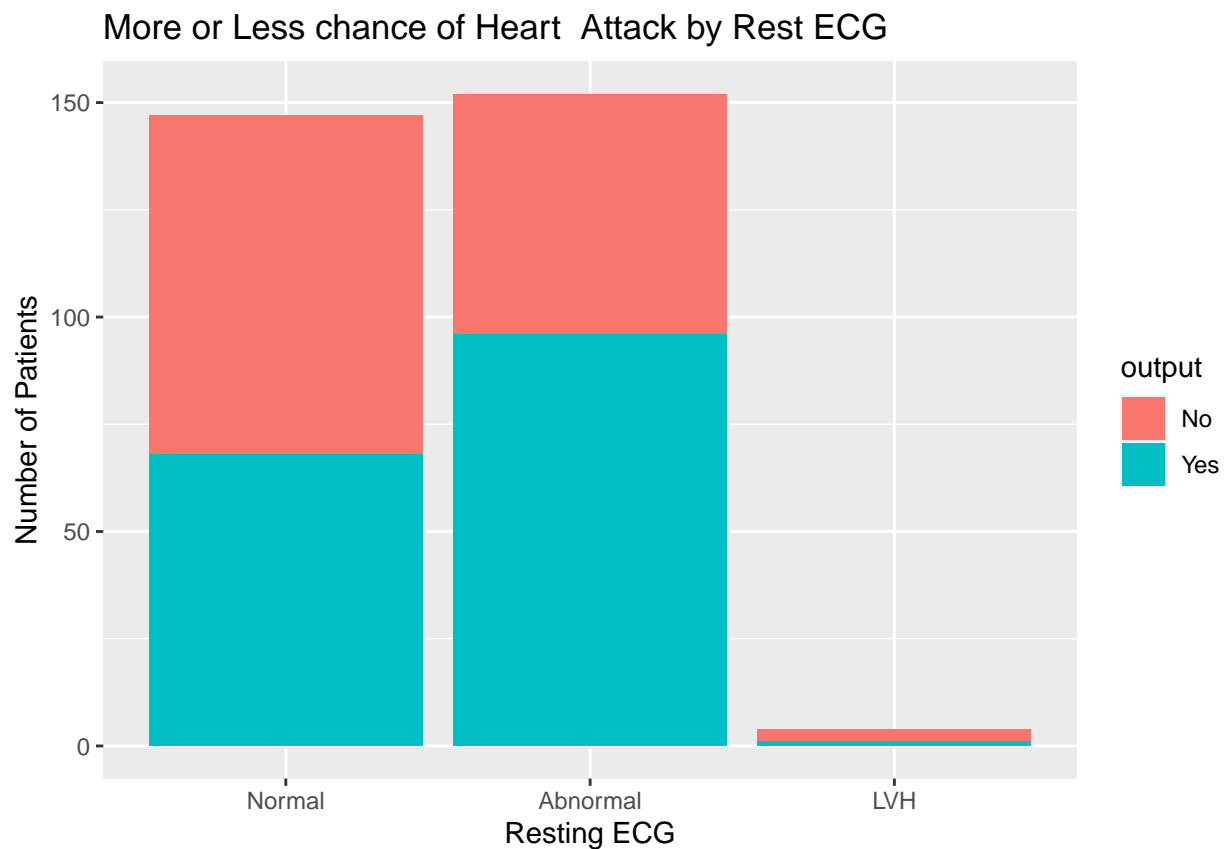
```
classofpar3 <- ggplot(heart, aes(x=exng, fill=output))+  
  geom_bar()+  
  xlab("Exercise Induced Agina")+  
  ylab("Number of Patients")+  
  ggtitle("More or Less chance of Heart Attack Exercise Induced Agina")  
classofpar3
```



```
classofpar4 <- ggplot(heart, aes(x=slp, fill=output))+  
  geom_bar()+  
  xlab("Slope of Peak Exercise")+  
  ylab("Number of Patients")+  
  ggtitle("More or Less chance of Heart Attack by Peak Exercise")  
classofpar4
```



```
classofpar5 <- ggplot(heart, aes(x=restecg, fill=output))+  
  geom_bar()+  
  xlab("Resting ECG")+  
  ylab("Number of Patients")+  
  ggtitle("More or Less chance of Heart Attack by Rest ECG")  
classofpar5
```




```
mfrow=par(c(2,3))
```

```
## Warning in par(c(2, 3)): argument 1 does not name a graphical parameter
```

```
library(ggplot2)
hist1 <- ggplot(heart, aes(x=age, fill=output))+
  geom_histogram()+
  xlab("Age")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack by Age")
hist1
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```


More or Less chance of Heart Attack by Age



```
hist22 <- ggplot(heart, aes(x=chol, fill=output))+
  geom_histogram()+
  xlab("Cholesterol")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack by cholesterol level")
hist22
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



More or Less chance of Heart Attack by cholesterol level



```
hist3 <- ggplot(heart, aes(x=trtbps, fill=output))+
  geom_histogram()+
  xlab("Resting BPS")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack by Resting BPS")
hist3
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

More or Less chance of Heart Attack by Resting BPS



```
hist4 <- ggplot(heart, aes(x=thalachh, fill=output))+
  geom_histogram()+
  xlab("Maximum Heart Rate")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack By Max Heart Rate")
hist4

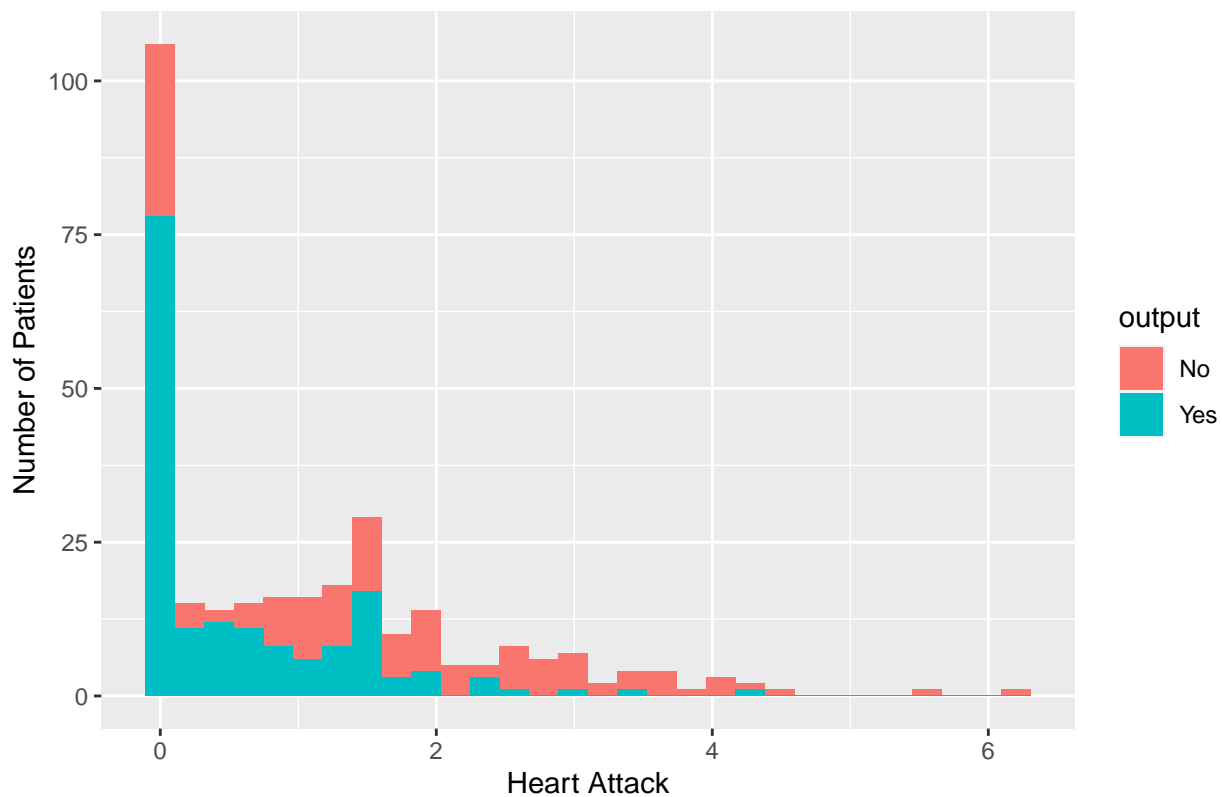
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
hist5 <- ggplot(heart, aes(x=oldpeak, fill=output))+
  geom_histogram()+
  xlab("Heart Attack")+
  ylab("Number of Patients")+
  ggtitle("More or Less chance of Heart Attack by ST Depression")
hist5
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

More or Less chance of Heart Attack by ST Depression



`summary(heart)`

```
##      age      sex      cp      trtbps
##  Min.   :29.00  Female: 96  typical   :143  Min.    : 94.0
##  1st Qu.:47.50  Male  :207  atypical   : 50  1st Qu.:120.0
##  Median :55.00                non-aginal : 87  Median :130.0
##  Mean   :54.37                Asymptomatic: 23  Mean   :131.6
##  3rd Qu.:61.00                                3rd Qu.:140.0
##  Max.   :77.00                                Max.   :200.0
##      chol      fbs      restecg      thalachh      exng
##  Min.   :126.0  False:258  Normal   :147  Min.    : 71.0  No :204
##  1st Qu.:211.0  True  : 45  Abnormal:152  1st Qu.:133.5  Yes: 99
##  Median :240.0                LVH      : 4  Median :153.0
##  Mean   :246.3                                Mean   :149.6
##  3rd Qu.:274.5                                3rd Qu.:166.0
##  Max.   :564.0                                Max.   :202.0
##      oldpeak      slp      caa      thall      output
##  Min.    :0.00  Upsloping : 21  Min.    :0.0000  Min.    :0.000  No :138
##  1st Qu.:0.00  Flat      :140  1st Qu.:0.0000  1st Qu.:2.000  Yes:165
##  Median :0.80  Downsloping:142  Median :0.0000  Median :2.000
##  Mean   :1.04                Mean   :0.7294  Mean   :2.314
##  3rd Qu.:1.60                3rd Qu.:1.0000  3rd Qu.:3.000
##  Max.   :6.20                Max.    :4.0000  Max.    :3.000
```

`str(heart)`

```
## 'data.frame': 303 obs. of 14 variables:
## $ age : int 63 37 41 56 57 57 56 44 52 57 ...
```

```
## $ sex      : Factor w/ 2 levels "Female","Male": 2 2 1 2 1 2 1 2 2 2 ...
## $ cp       : Factor w/ 4 levels "typical","atypical",...: 4 3 2 2 1 1 2 2 3 3 ...
## $ trtbps   : int   145 130 130 120 120 140 140 120 172 150 ...
## $ chol     : int   233 250 204 236 354 192 294 263 199 168 ...
## $ fbs      : Factor w/ 2 levels "False","True": 2 1 1 1 1 1 1 1 2 1 ...
## $ restecg  : Factor w/ 3 levels "Normal","Abnormal",...: 1 2 1 2 2 2 1 2 2 2 ...
## $ thalachh : int   150 187 172 178 163 148 153 173 162 174 ...
## $ exng     : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 1 1 1 1 1 ...
## $ oldpeak  : num    2.3 3.5 1.4 0.8 0.6 0.4 1.3 0 0.5 1.6 ...
## $ slp      : Factor w/ 3 levels "Upsloping","Flat",...: 1 1 3 3 3 2 2 3 3 3 ...
## $ caa      : int    0 0 0 0 0 0 0 0 0 0 ...
## $ thall    : int    1 2 2 2 2 1 2 3 3 2 ...
## $ output   : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
```

Assumptions 1. Normality 2. Positive Determinant for variance co-variance matrix 3. Equal variance between two groups

“(d) Providing univariate means and variances

```
df <- data.frame(heart$age,heart$trtbps,heart$chol,heart$thalachh,heart$oldpeak)
by(df, heart$output, colMeans)
```

```
## heart$output: No
##      heart.age  heart.trtbps      heart.chol heart.thalachh  heart.oldpeak
##      56.601449   134.398551    251.086957    139.101449      1.585507
## -----
## heart$output: Yes
##      heart.age  heart.trtbps      heart.chol heart.thalachh  heart.oldpeak
##      52.4969697  129.3030303    242.2303030    158.4666667      0.5830303
```

```
by(df, heart$output, var)
```

```
## heart$output: No
##      heart.age  heart.trtbps      heart.chol heart.thalachh
## heart.age      63.394742    35.831535    44.3414789    -23.886279
## heart.trtbps    35.831535    350.810801    125.3592510     2.368031
## heart.chol      44.341479    125.359251    2445.7588067     50.297683
## heart.thalachh -23.886279     2.368031    50.2976833     510.704961
## heart.oldpeak   1.066445     3.128446    -0.6016503     -6.267132
##
##      heart.oldpeak
## heart.age      1.066445
## heart.trtbps    3.128446
## heart.chol      -0.6016503
## heart.thalachh  -6.2671321
## heart.oldpeak   1.6908833
## -----
## heart$output: Yes
##      heart.age  heart.trtbps      heart.chol heart.thalachh  heart.oldpeak
## heart.age      91.214930    42.421656    131.525092    -96.288211     1.3017775
## heart.trtbps    42.421656    261.456393     80.783444     8.693089     2.2911493
## heart.chol      131.525092     80.783444    2867.910052    14.843089     2.4130783
## heart.thalachh -96.288211     8.693089    14.843089    367.652846    -2.7249593
## heart.oldpeak   1.301778     2.291149     2.413078    -2.724959     0.6094664
```

```
by(df, heart$output, cor)
```

```
## heart$output: No
```

```
##          heart.age heart.trtbps heart.chol heart.thalachh heart.oldpeak
## heart.age      1.0000000  0.24027153  0.11260994   -0.13275071   0.10300417
## heart.trtbps    0.2402715  1.00000000  0.13533593    0.00559456   0.12845037
## heart.chol      0.1126099  0.13533593  1.00000000    0.04500452  -0.00935579
## heart.thalachh -0.1327507  0.00559456  0.04500452    1.00000000  -0.21326864
## heart.oldpeak   0.1030042  0.12845037 -0.00935579   -0.21326864   1.00000000
```

```
## -----
## heart$output: Yes
##          heart.age heart.trtbps heart.chol heart.thalachh heart.oldpeak
## heart.age      1.0000000  0.27469770  0.25715377   -0.52580074   0.17459385
## heart.trtbps    0.2746977  1.00000000  0.09329105    0.02803855   0.18150094
## heart.chol      0.2571538  0.09329105  1.00000000    0.01445515   0.05771833
## heart.thalachh -0.5258007  0.02803855  0.01445515    1.00000000  -0.18203973
## heart.oldpeak   0.1745938  0.18150094  0.05771833   -0.18203973   1.00000000
```

```
library(mvnormtest)
multv <- t(df)

mshapiro.test(multv)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  Z
## W = 0.93989, p-value = 9.335e-10
```

Highly significant P value we might have to reject the null hypothesis, does not pass normality test.

```
covheart <- cov(df, method = 'spearman')
det(covheart)
```

```
## [1] 1.476375e+19
```


It passes the positive determinant test of variance-covariance matrix.

```
#Graph the means of the 5 variables for presence or absence of heart attack
library(gplots)
```

```
##
## Attaching package: 'gplots'
## The following object is masked from 'package:stats':
##
##      lowess
```

```
plotmeans(heart$age~heart$output, data=heart,ylim=c(0,100),xlab="Heart Attack",legends = c("No","Yes"),
```


likelihood of Heart Attack



Heart Attack

```
plotmeans(heart$strtbps~heart$output, data=heart,ylim=c(0,200),xlab="Heart Attack",legends = c("No","Yes"))
```


likelihood of Heart Attack



Heart Attack

```
plotmeans(heart$chol~heart$output, data=heart,ylim=c(0,300),xlab="Heart Attack",legends = c("No","Yes"))
```


likelihood of Heart Attack



Heart Attack

```
plotmeans(heart$thalachh~heart$output, data=heart,ylim=c(0,200),xlab="Heart Attack",legends = c("No","Yes"))
```

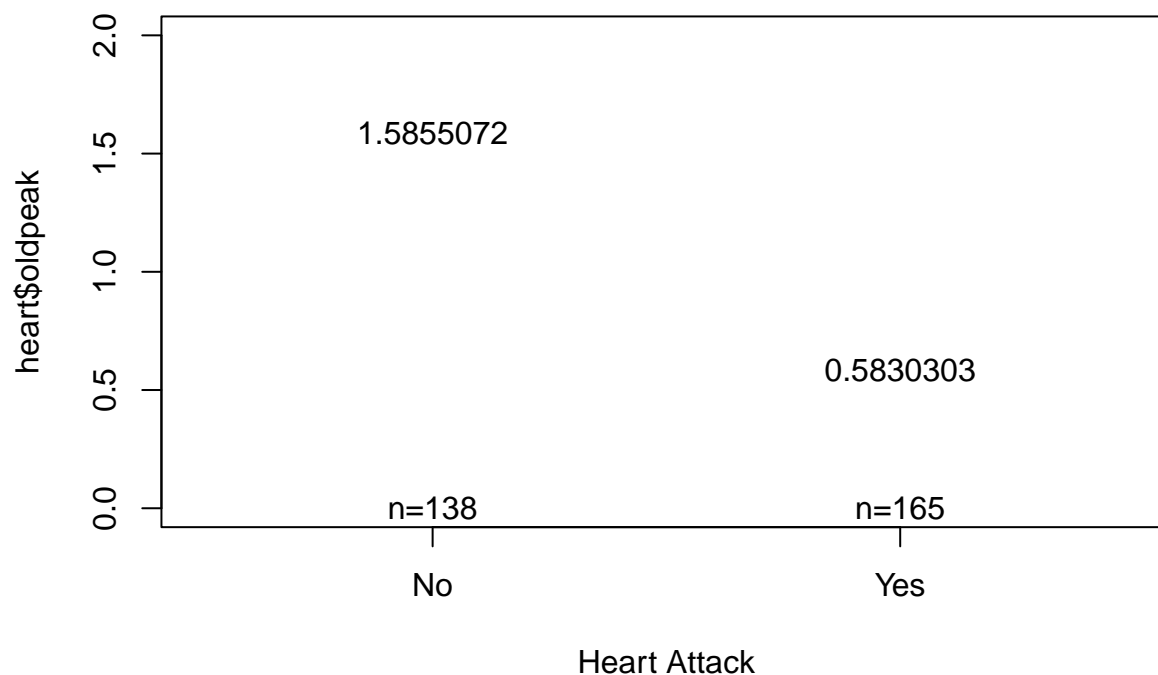
likelihood of Heart Attack



Heart Attack

```
plotmeans(heart$oldpeak~heart$output, data=heart,ylim=c(0,2),xlab="Heart Attack",legends = c("No","Yes"))
```


likelihood of Heart Attack



```
library(ICSNP)
```

```
## Loading required package: mvtnorm
```

```
## Loading required package: ICS
```

```
library(mvtnorm)
```

```
result <- HotellingsT2(df[1:165,], df[166:303,])
```

```
# Print the test result
```

```
print(result)
```

```
##
```

```
## Hotelling's two sample T2-test
```

```
##
```

```
## data: df[1:165, ] and df[166:303, ]
```

```
## T.2 = 22.924, df1 = 5, df2 = 297, p-value < 2.2e-16
```

```
## alternative hypothesis: true location difference is not equal to c(0,0,0,0,0)
```

```
library(caTools)
```

```
set.seed(1028)
```

```
total_rows <- nrow(heart)
```

```
# Calculate the number of rows for the training and testing sets
```

```
train_rows <- round(0.8 * total_rows) # 80% for training
```

```
test_rows <- total_rows - train_rows # 20% for testing
```

```
# Generate random indices for the training set
```

```
train_indices <- sample(1:total_rows, train_rows)
```

```
# Create the training and testing datasets
```

```
train_data <- heart[train_indices, ]
test_data <- heart[-train_indices, ]
```

```
set.seed(1130)
dim(train_data)
```

```
## [1] 242 14
```

```
dim(test_data)
```

```
## [1] 61 14
```

- (d) Perform Random Forest on the training data in order to predict origin2. Evaluate performance on the test data. What test error do you obtain?

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
set.seed(937)
```

```
#random forest model
```

```
randori <- randomForest(output~.-output, train_data, mtry =5, importance = TRUE)
```

```
#predicting new data on the random forest model
```

```
pred.randori <- predict(randori,newdata = test_data)
```

```
mean(pred.randori!=test_data$output)
```

```
## [1] 0.147541
```

```
mean(pred.randori==test_data$output)
```

```
## [1] 0.852459
```

```
#confusion matrix
```

```
table(pred.randori,test_data$output)
```

```
##
```

```
## pred.randori No Yes
```

```
##           No  21   2
```

```
##           Yes  7  31
```

```
#importance of variables
```


```
importance(randori)
```

##		No	Yes	MeanDecreaseAccuracy	MeanDecreaseGini
##	age	4.3857870	11.4440591	11.4348948	11.6390449
##	sex	6.4744027	10.5151711	11.8072660	3.4463238
##	cp	19.2109285	15.4486309	23.3534498	18.5684786
##	trtbps	0.1017973	3.7541518	2.8876303	9.4932074
##	chol	-2.1180602	0.7002795	-0.8142109	9.2332600
##	fbs	-2.3740169	-0.3350165	-1.7114196	0.7235888
##	restecg	-0.9285768	-0.1308484	-0.8094133	1.4471774

```
## thalachh 4.7758225 7.2811953 8.8259248 13.6352038
## exng 9.0476382 3.6010811 8.4185208 6.1034752
## oldpeak 14.2753390 11.8247450 19.0113823 14.0534116
## slp 6.3329648 4.1100172 7.6159437 4.5372944
## caa 15.0044316 18.9716049 22.6166979 14.4132533
## thall 9.9129996 15.7886058 17.2388402 12.0861334
```

```
#plot of variables based on importance
varImpPlot(randori)
```

randori



These plots above show important variables and cp variable shows up as the most important variable.

- (c) Perform Classification Tree on the training data in order to predict output. Use cross- validation to prune the tree. Plot the resulting tree. Evaluate performance on the test data. What test error do you obtain?

```
library(tree)
attach(heart)
#change this to factors so that it can make a classification tree without having level issues
output= factor(output)
tree_output <- tree(output~.-output,train_data)

summary(tree_output)#gives summary statistics of tree
```

```
##
## Classification tree:
## tree(formula = output ~ . - output, data = train_data)
## Variables actually used in tree construction:
## [1] "cp"      "caa"      "oldpeak" "thall"    "thalachh" "chol"    "age"
## [8] "sex"
## Number of terminal nodes: 19
```

```

## Residual mean deviance:  0.4743 = 105.8 / 223
## Misclassification error rate: 0.1116 = 27 / 242

set.seed(4399)
#perform prediction on test data
tree_pred = predict(tree_output, test_data, type = "class")

table(tree_pred, test_data$output) #confusion matrix

##
## tree_pred No Yes
##      No  22  4
##      Yes  6 29

mean(tree_pred != test_data$output) #test error

## [1] 0.1639344

mean(tree_pred == test_data$output) #test accuracy

## [1] 0.8360656


set.seed(1111)
#perform cross validation
cv_output = cv.tree(tree_output, FUN = prune.misclass)

cv_output

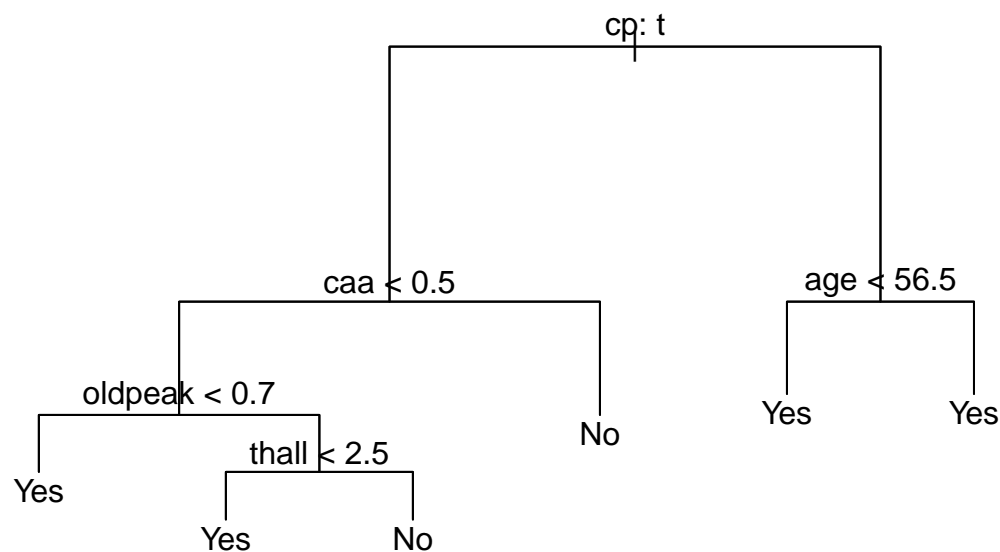
## $size
## [1] 19 13  8  6  4  2  1
##
## $dev
## [1]  50  49  44  53  55  60 112
##
## $k
## [1] -Inf    0    1    3    4    7   50
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune"          "tree.sequence"

#plot the CV
plot(cv_output$size, cv_output$dev, type = "b")
text(tree_output, pretty=TRUE, cex=0.8)

```



```
#prune tree and plot resulting tree
prune.output <- prune.tree(tree_output, best = 6)
plot(prune.output)
text(prune.output, pretty=TRUE)
```



find below the test terror for my pruned classification tree

```
#predict test data on pruned tree
prune.pred = predict(prune.output, test_data, type="class")
table(prune.pred, test_data$output)
```

```
##
## prune.pred No Yes
##      No  20   0
##      Yes   8  33
```

```
mean(prune.pred != test_data$output) #test error
```

```
## [1] 0.1311475
```

```
mean(prune.pred == test_data$output) #test accuracy for pruned tree
```

```
## [1] 0.8688525
```

- e) Fit a Support Vector Classifier to the data with various values of cost, in order to predict the likelihood of a heart attack or not . Report the cross-validation errors associated with different values of this parameter. Comment on your results.

```
library(e1071)
```

```
set.seed(821)
```

```
# SV classifier model
```

```
svmlinear = svm(output ~ .-output, data = train_data, kernel = "linear", ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)), summary(svmlinear))
```

```
##
```

```
## Call:
```

```
## svm(formula = output ~ . - output, data = train_data, kernel = "linear",  
##      ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
```

```
##
```

```
##
```

```
## Parameters:
```

```
##   SVM-Type:  C-classification
```

```
##   SVM-Kernel: linear
```

```
##         cost:  1
```

```
##
```

```
## Number of Support Vectors:  107
```

```
##
```

```
## ( 52 55 )
```

```
##
```

```
##
```

```
## Number of Classes:  2
```

```
##
```

```
## Levels:
```

```
## No Yes
```

```
#prediction of SV classifier, table and error
```

```
pred.svmlinear <- predict(svmlinear,newdata = test_data)
```

```
table(test_data$output,pred.svmlinear)
```

```
##      pred.svmlinear
```

```
##      No Yes
```

```
## No  21  7
```

```
## Yes  1 32
```

```
mean(pred.svmlinear!=test_data$output)#test error
```

```
## [1] 0.1311475
```

```
mean(pred.svmlinear==test_data$output)#test accuracy
```

```
## [1] 0.8688525
```

As can be seen above the cross-validation gives a better error than just the classifier model

```

set.seed(433)
#cv on SV classifier
svmlineart = tune(svm,output ~ .-output, data = train_data, kernel = "linear", ranges = list(cost = c(
summary(svmlineart)

```

```

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   100
##
## - best performance: 0.1946667
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.4545000 0.06114166
## 2 1e-02 0.2151667 0.05860229
## 3 1e+00 0.1986667 0.07286805
## 4 5e+00 0.2030000 0.08022268
## 5 1e+01 0.1988333 0.07857547
## 6 1e+02 0.1946667 0.07664170

```

The lowest error was achieved by using a cost of 100, which was 0.194. The highest error was 0.4545 which was when cost was set to 0.001.

- f) Now repeat (e), this time using Support Vector Machines (SVMs) with radial and polynomial basis kernels, with different values of gamma and degree and cost. Comment on your results.

```

set.seed(439)
#SVM polynomial model
svmpol = svm(output ~ .-output, data = train_data, kernel = "polynomial", ranges = list(cost = c(0.1, 1, 5, 10), degree = c(2, 3, 4)))
summary(svmpol)

```

```

##
## Call:
## svm(formula = output ~ . - output, data = train_data, kernel = "polynomial",
##   ranges = list(cost = c(0.1, 1, 5, 10), degree = c(2, 3, 4)))
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##       cost:  1
##     degree:  3
##    coef.0:  0
##
## Number of Support Vectors:  192
##
##   ( 95 97 )
##
##
## Number of Classes:  2

```

```
##
## Levels:
## No Yes
#predicting new data on the SVM model polynomial model
pred.svm.pol <- predict(svm.pol, newdata = test_data)
#confusion matrix
table(test_data$output, pred.svm.pol)

##      pred.svm.pol
##      No Yes
## No  16  12
## Yes  0  33

#computing test error obtained
mean(pred.svm.pol != test_data$output)

## [1] 0.1967213

#computing test Accuracy obtained
mean(pred.svm.pol == test_data$output)

## [1] 0.8032787

set.seed(1145)
svmpolt = tune(svm, output ~ .-output, data = train_data, kernel = "polynomial", ranges = list(cost = c(0.1, 1, 5, 10), degree = c(2, 3, 4)))
summary(svmpolt)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##    10      3
##
## - best performance: 0.2021667
##
## - Detailed performance results:
##   cost degree      error dispersion
## 1  0.1      2 0.4298333 0.08117748
## 2  1.0      2 0.2025000 0.07964524
## 3  5.0      2 0.2105000 0.07795238
## 4 10.0      2 0.2228333 0.07521184
## 5  0.1      3 0.4381667 0.08838555
## 6  1.0      3 0.2855000 0.05503114
## 7  5.0      3 0.2148333 0.06421150
## 8 10.0      3 0.2021667 0.07073272
## 9  0.1      4 0.4548333 0.09699322
##10  1.0      4 0.3885000 0.07872224
##11  5.0      4 0.2938333 0.07067160
##12 10.0      4 0.3100000 0.09032178
```

The lowest error was achieved by using a cost of 1 and degree of 2, which was 0.2025. The highest error was 0.4548 which was when cost was set to 0.1 and a degree 4.


```

set.seed(1146)
svmrad = svm(output ~ .-output, data = train_data, kernel = "radial", ranges = list(cost = c(0.1, 1, 5),
summary(svmrad)

##
## Call:
## svm(formula = output ~ . - output, data = train_data, kernel = "radial",
##      ranges = list(cost = c(0.1, 1, 5, 10), gamma = c(0.01, 0.1, 1,
##      5, 10, 100)))
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:   1
##
## Number of Support Vectors: 139
##
## ( 70 69 )
##
##
## Number of Classes: 2
##
## Levels:
## No Yes
#prediction of SVM radial classifier, table and error
pred.svmrad <- predict(svmrad,newdata = test_data)
table(test_data$output,pred.svmrad)

##      pred.svmrad
##      No Yes
## No  24  4
## Yes  1 32
mean(pred.svmrad!=test_data$output)

## [1] 0.08196721
mean(pred.svmrad==test_data$output)

## [1] 0.9180328
set.seed(934)
svmradt = tune(svm,output ~ .-output, data = train_data, kernel = "radial", ranges = list(cost = c(0.1
summary(svmradt)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   10 0.01
##
## - best performance: 0.2061667

```

```
##
## - Detailed performance results:
##      cost gamma      error dispersion
## 1   0.1 1e-02 0.3760000 0.17869935
## 2   1.0 1e-02 0.2185000 0.07370884
## 3   5.0 1e-02 0.2103333 0.05470347
## 4  10.0 1e-02 0.2061667 0.06543327
## 5   0.1 1e-01 0.2473333 0.09171278
## 6   1.0 1e-01 0.2145000 0.05944211
## 7   5.0 1e-01 0.2271667 0.05885245
## 8  10.0 1e-01 0.2556667 0.08875059
## 9   0.1 1e+00 0.4546667 0.14808614
## 10  1.0 1e+00 0.3926667 0.14146980
## 11  5.0 1e+00 0.3926667 0.12451932
## 12 10.0 1e+00 0.3926667 0.12451932
## 13  0.1 5e+00 0.4546667 0.14808614
## 14  1.0 5e+00 0.4546667 0.14808614
## 15  5.0 5e+00 0.4546667 0.14808614
## 16 10.0 5e+00 0.4546667 0.14808614
## 17  0.1 1e+01 0.4546667 0.14808614
## 18  1.0 1e+01 0.4546667 0.14808614
## 19  5.0 1e+01 0.4546667 0.14808614
## 20 10.0 1e+01 0.4546667 0.14808614
## 21  0.1 1e+02 0.4546667 0.14808614
## 22  1.0 1e+02 0.4546667 0.14808614
## 23  5.0 1e+02 0.4546667 0.14808614
## 24 10.0 1e+02 0.4546667 0.14808614
```

The lowest error was achieved by using a cost of 5 and gamma of 0.01, which was 0.2103. The highest error was 0.4547 which was when cost was set to 0.1 and a gamma of 1. And the lowest CV error is not lower than that of the regular model which was not expected.

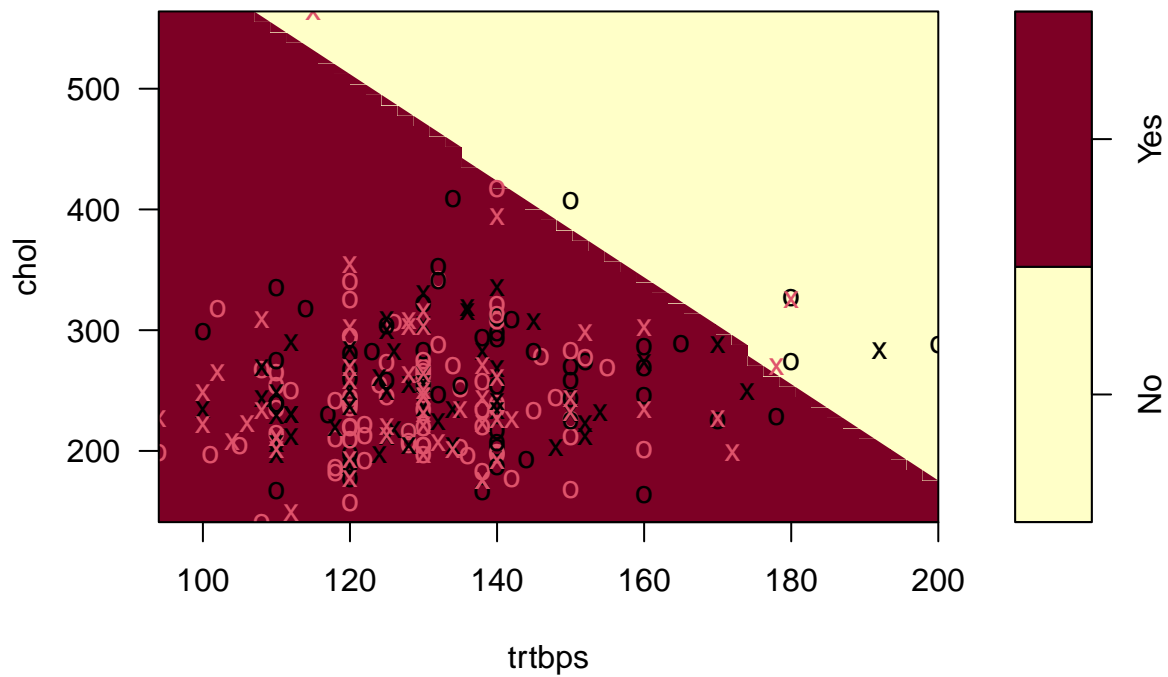
(g) Make some plots to back up your assertions in (e) and (f).

```
library(kernlab)
```

```
##
## Attaching package: 'kernlab'


## The following object is masked from 'package:ggplot2':
##
##      alpha
plot(svmlinear,train_data,col~trtbps)
```

SVM classification plot




```
plot(svmlinear,train_data,trtbps~age)
```

SVM classification plot



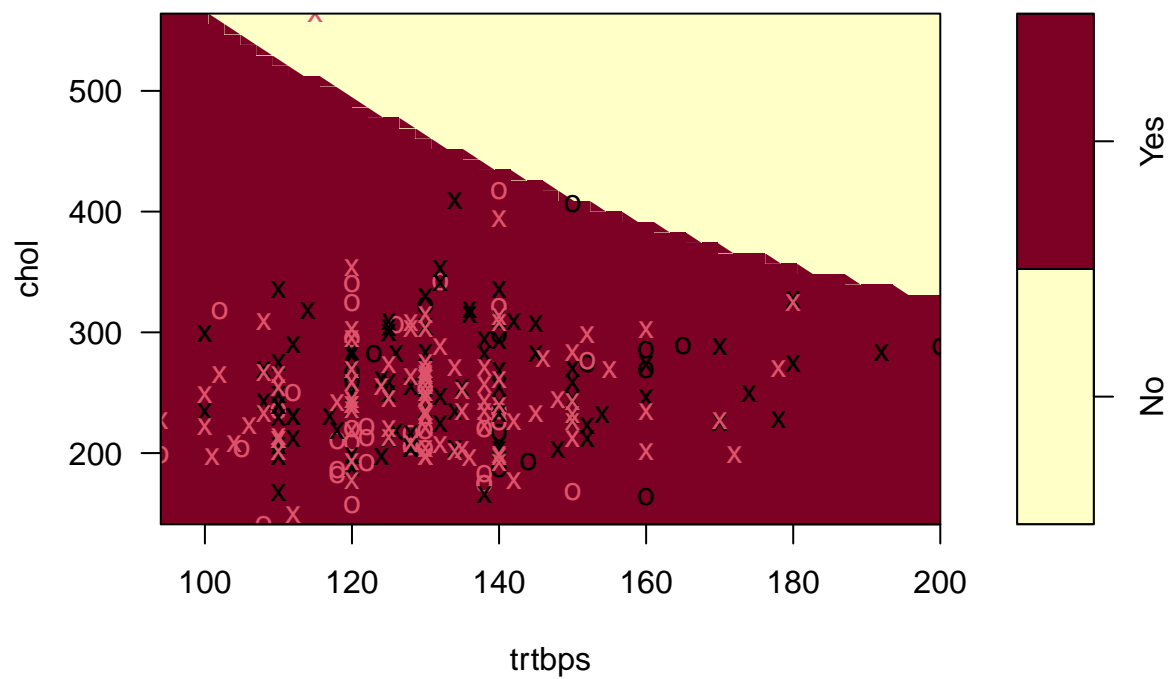
```
plot(svmlinear,train_data,thalachh~trtbps)
```

SVM classification plot



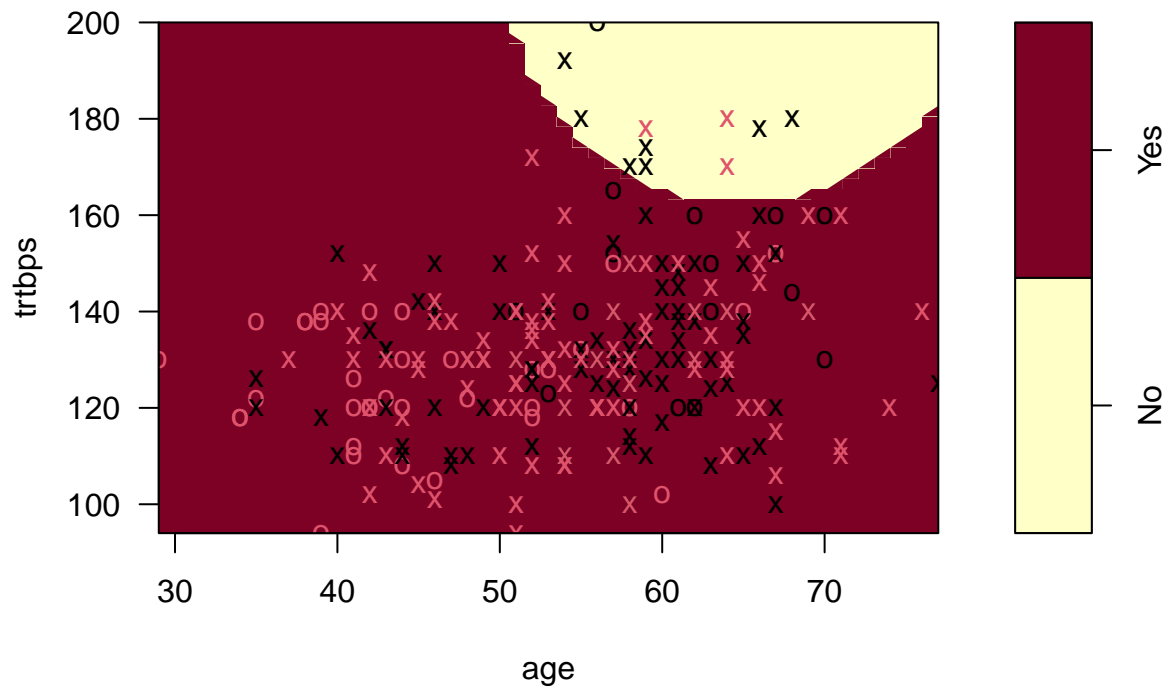
```
plot(svm.pol, train_data, chol ~ trtbps)
```

SVM classification plot



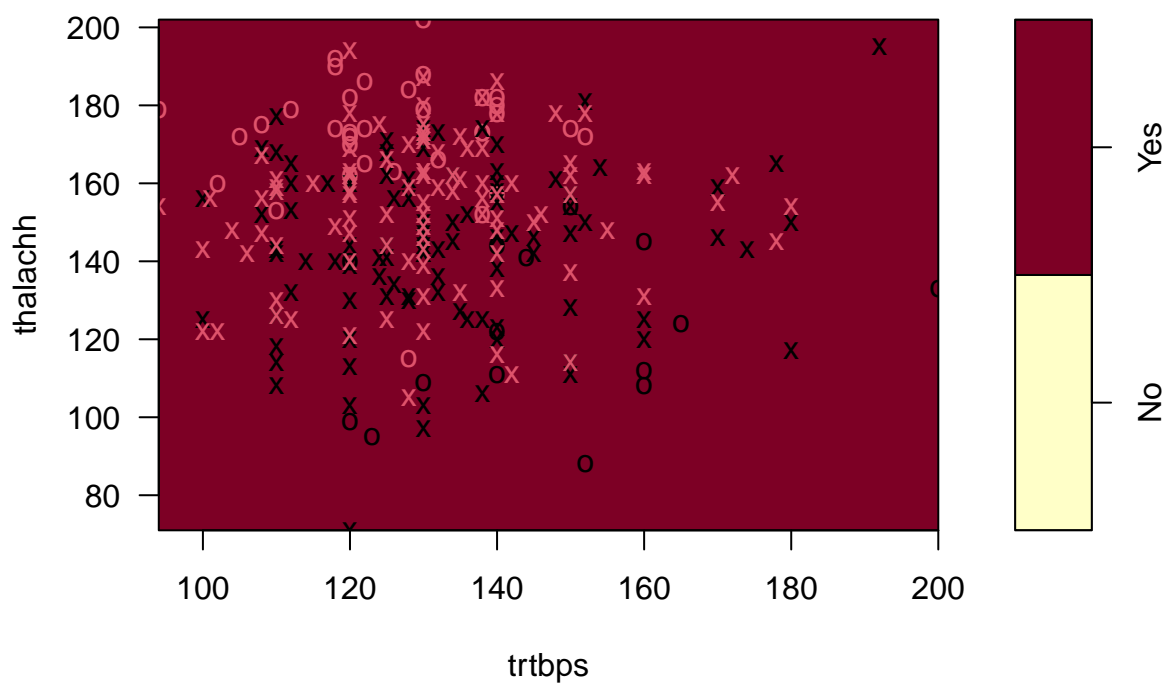
```
plot(svm.pol, train_data, trtbps ~ age)
```

SVM classification plot




```
plot(svmrad,train_data,thalachh~trtbps)
```

SVM classification plot




```
plot(svmrad,train_data,chol~trtbps)
```

SVM classification plot



```
plot(svmrad,train_data,trtbps~age)
```

SVM classification plot



```
plot(svmrad,train_data,age~thalachh)
```

SVM classification plot

